

# Real-World CKA



Domains	Description
Cluster Architecture, Installation, and Configuration	All things infrastructure. How does Kubernetes ACTUALLY work?
Workload Scheduling	Figuring out how Kubernetes has the ability to run workloads
Services and Networking	Kubernetes is nothing without networking. It's how the entire platform works.
Storage	Think hard drives or storage that isn't ephemeral
Troubleshooting	Understanding how to fix issues from the cluster to the workloads to the networks.
The Real World	A cert cannot fully prepare you for the real world. Let's dive into what you need.

# CKA vs CKAD

---

---

- Infra focused
  - Environment configurations at the platform level
- Developer focused
  - Environment configuration at the application level

# CKA vs KCNA

---

---

- Infra focused
- Environment configurations at the platform level
- Same focus as the CKA, just at a lower level

# CKA vs CKS

---

---

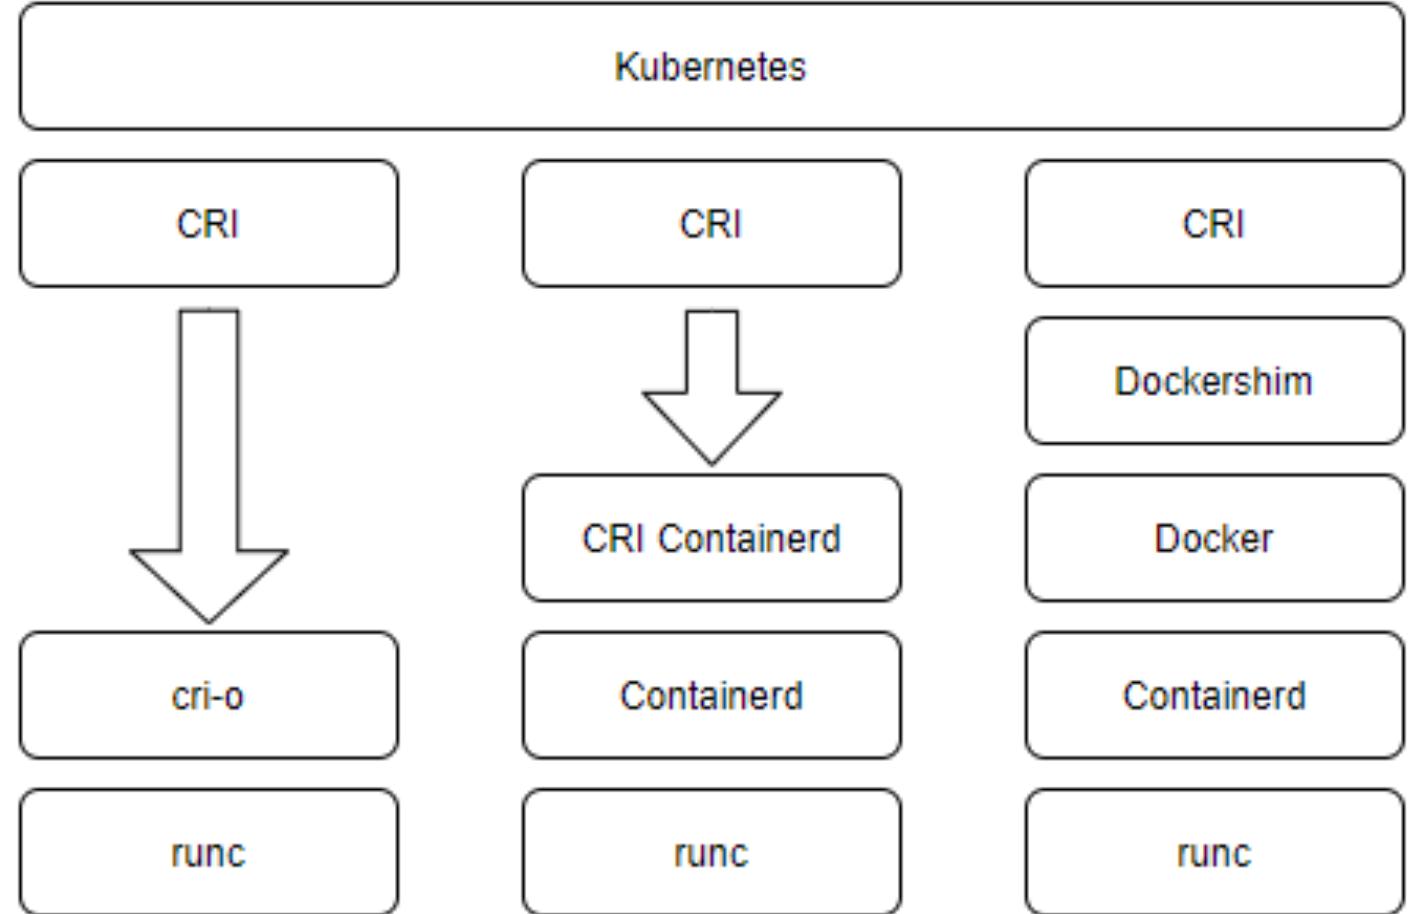
- Infra focused
  - Environment configurations at the platform level
- 
- Everything and anything at the security level

# CRI-O Runtime

- CRI Compliant
- Solid out-of-the-box implementation for Kubeadm
- Can be used in the cloud



# CRI-O Runtime

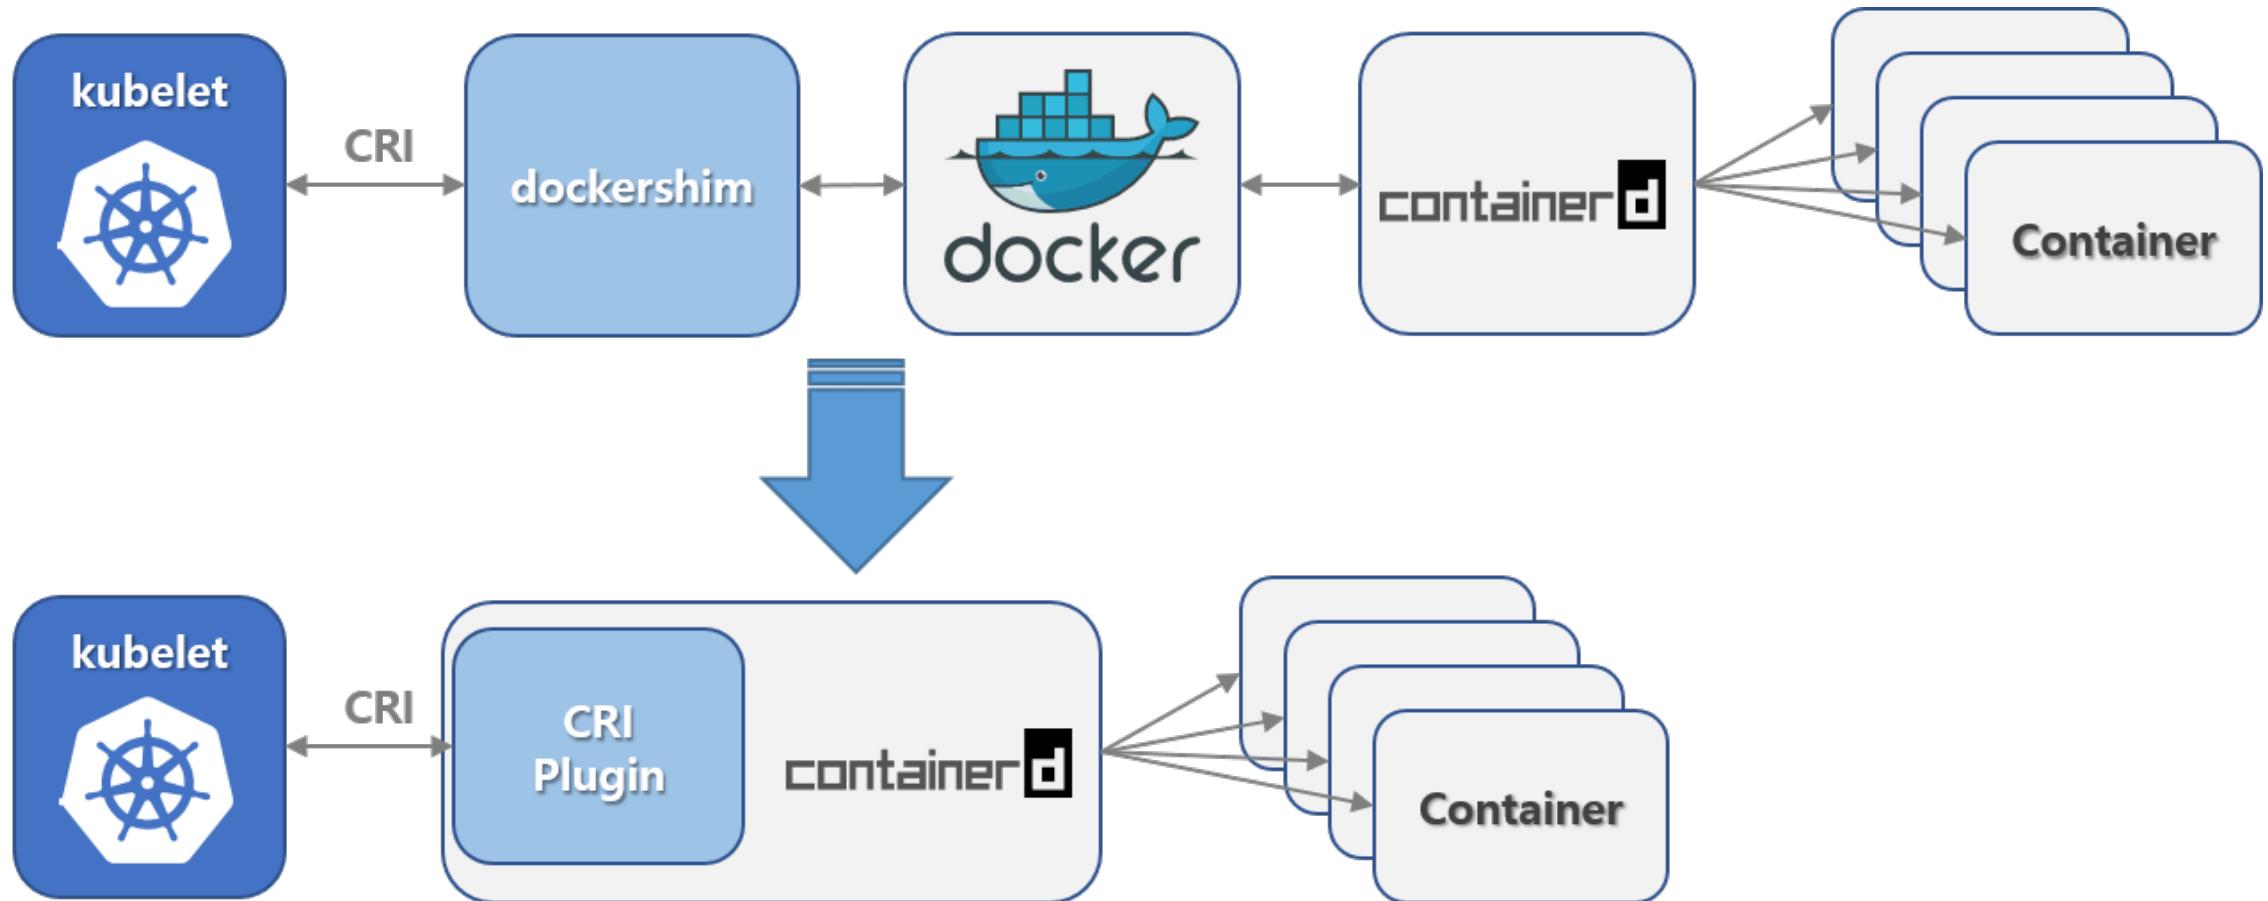


# Containerd Runtime

- Graduated CNCF Project
- Available as a Daemon for Windows and Linux
  - OCI Compliant
  - Default for AKS and EKS



# Containerd Runtime



# Kubernetes Version Information

- Kubernetes versions are expressed as x.y.z, where x is the major version, y is the minor version, and z is the patch version, following Semantic Versioning terminology.
- The Kubernetes project maintains release branches for the most recent three minor releases (For example - 1.28, 1.27, 1.26).
- Kubernetes 1.19 and newer receive approximately 1 year of patch support.
- Kubernetes 1.18 and older received approximately 9 months of patch support.

# Versioning

---

KUBERNETES “SEMANTIC” VERSIONING EXPLAINED



# DNS

---

- The Domain Name System is a hierarchical and distributed naming system for computers, services, and other resources in the Internet or other Internet Protocol networks.



# CoreDNS

---

---

The Kubernetes default

# eBPF



Remove the need for ip-tables, which means better scalability



Far better kernel efficiency as it's JIT compiled and runs directly in the kernel



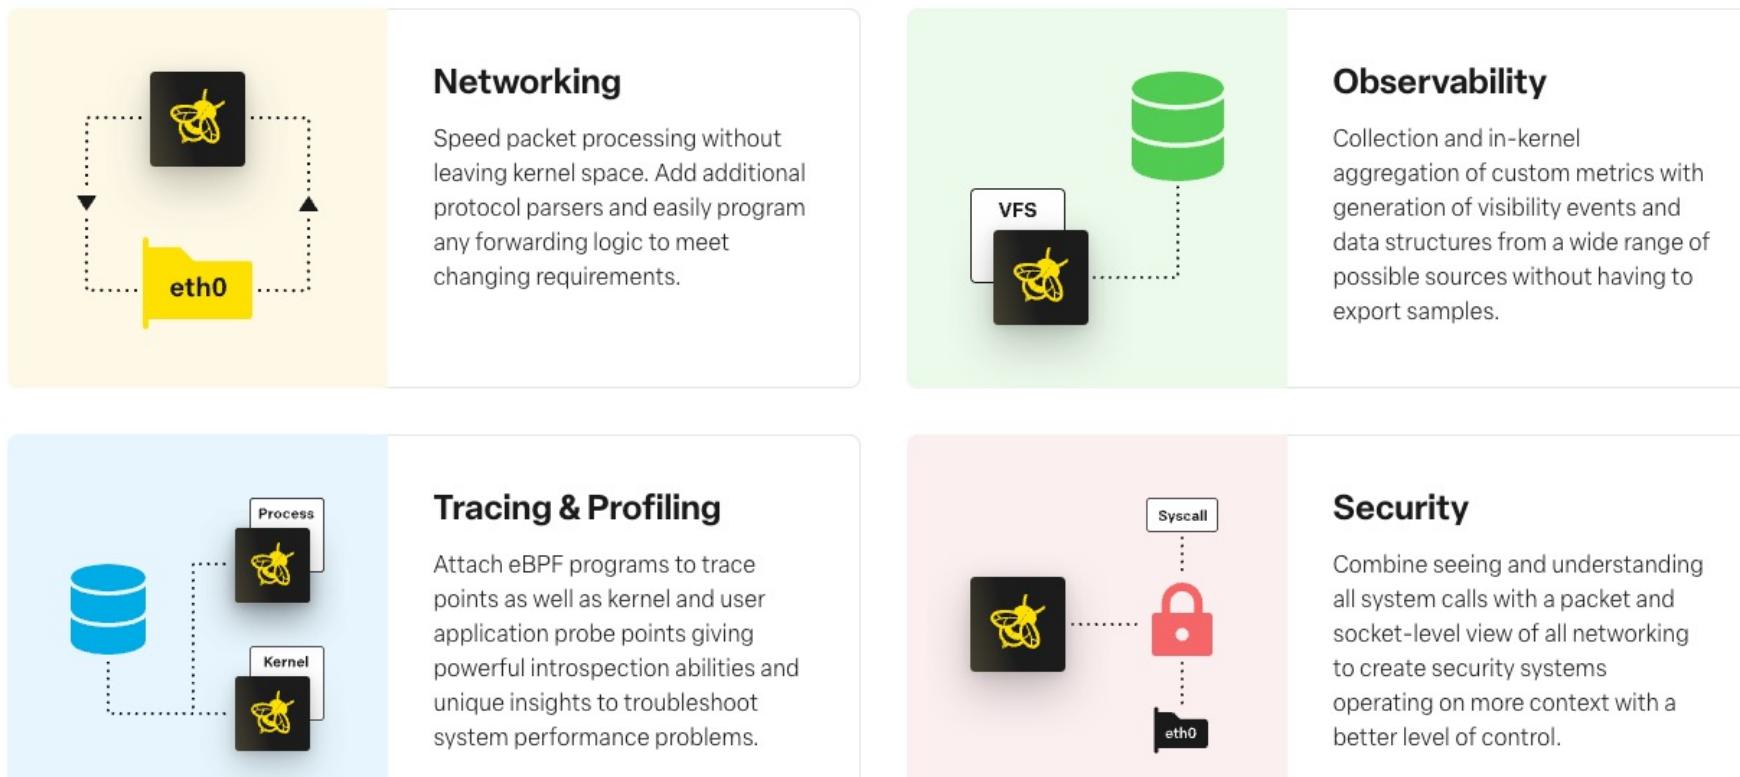
Networking, observability, tracing, and security out of the box



Used by organizations such as Google and Netflix

# eBPF

## What's possible with eBPF?

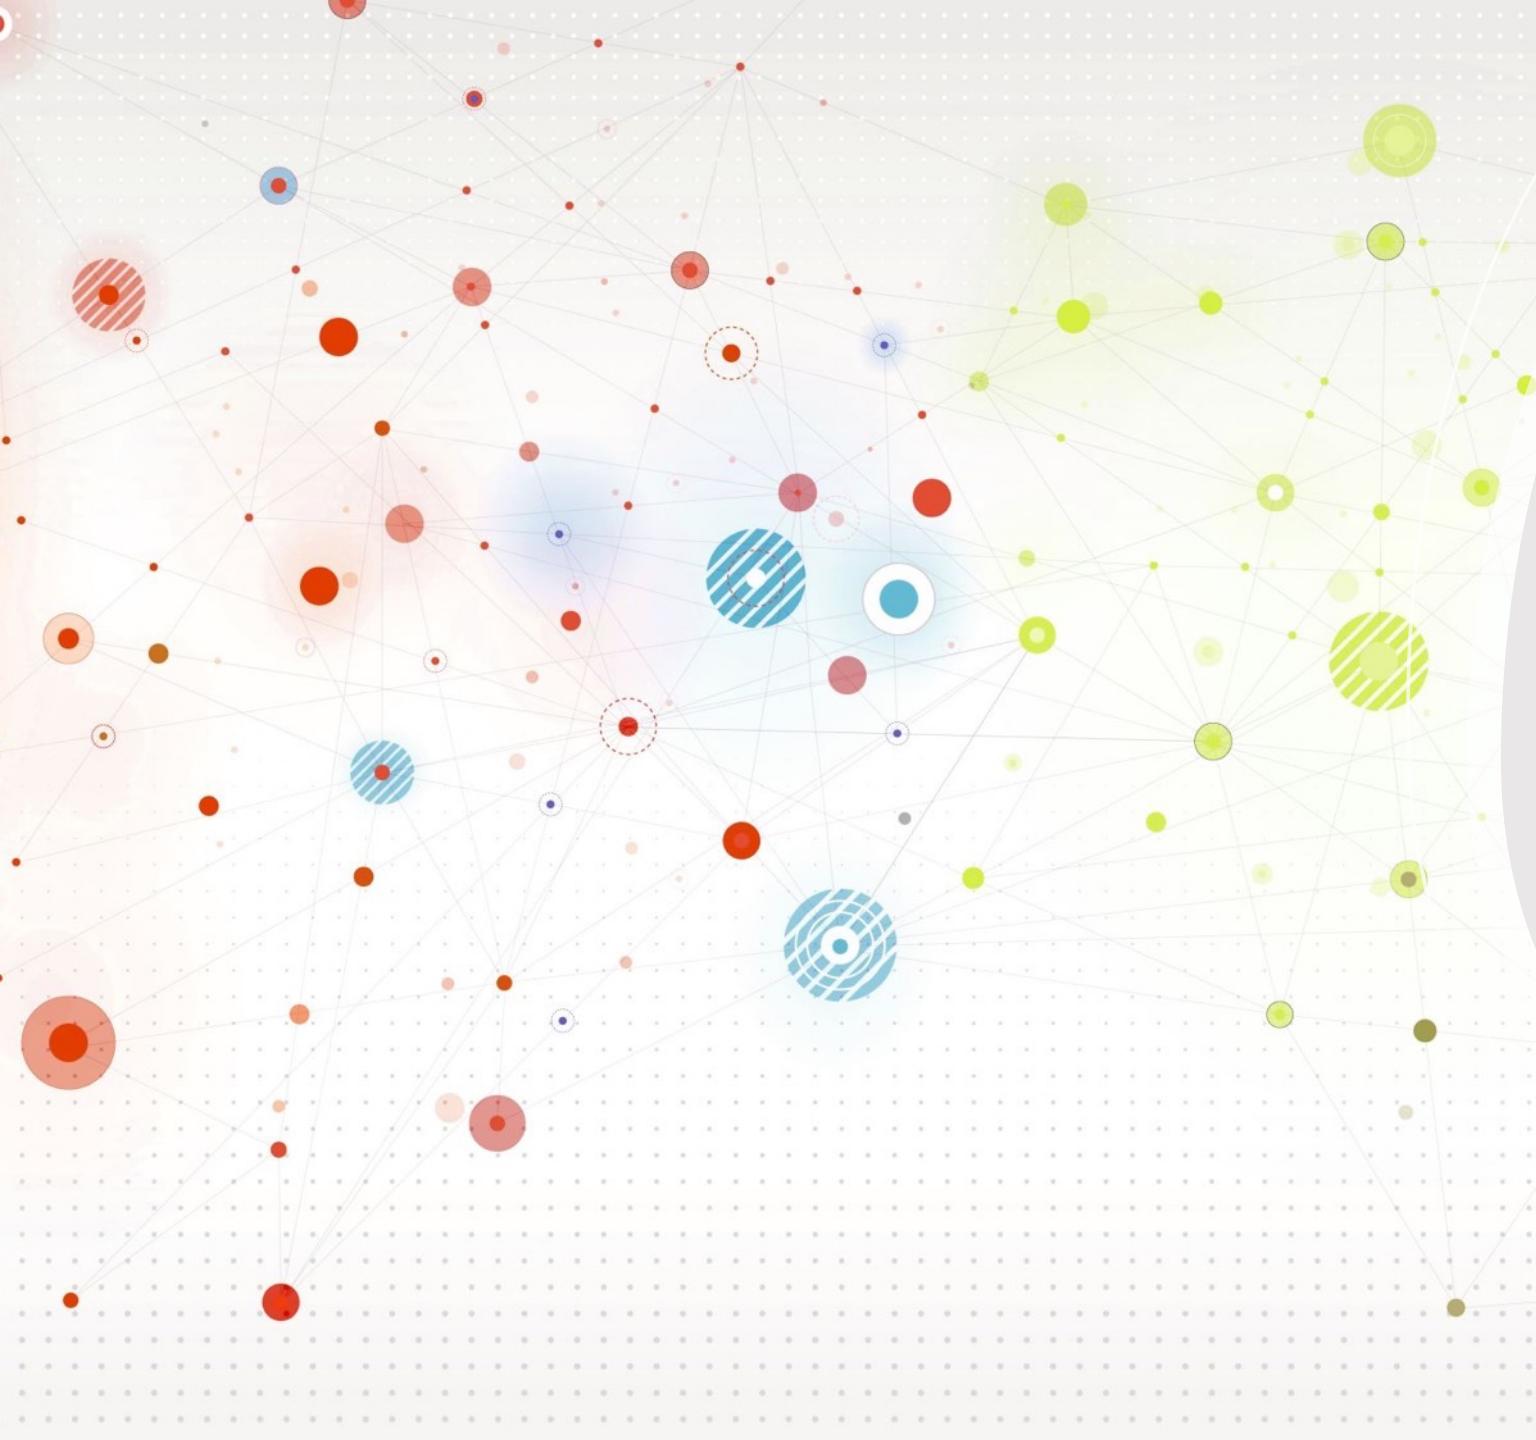


Source: <https://ebpf.io/static/observability-0e7ded4a90e87ea6a804f95a26f0ca71.svg>

# Pods

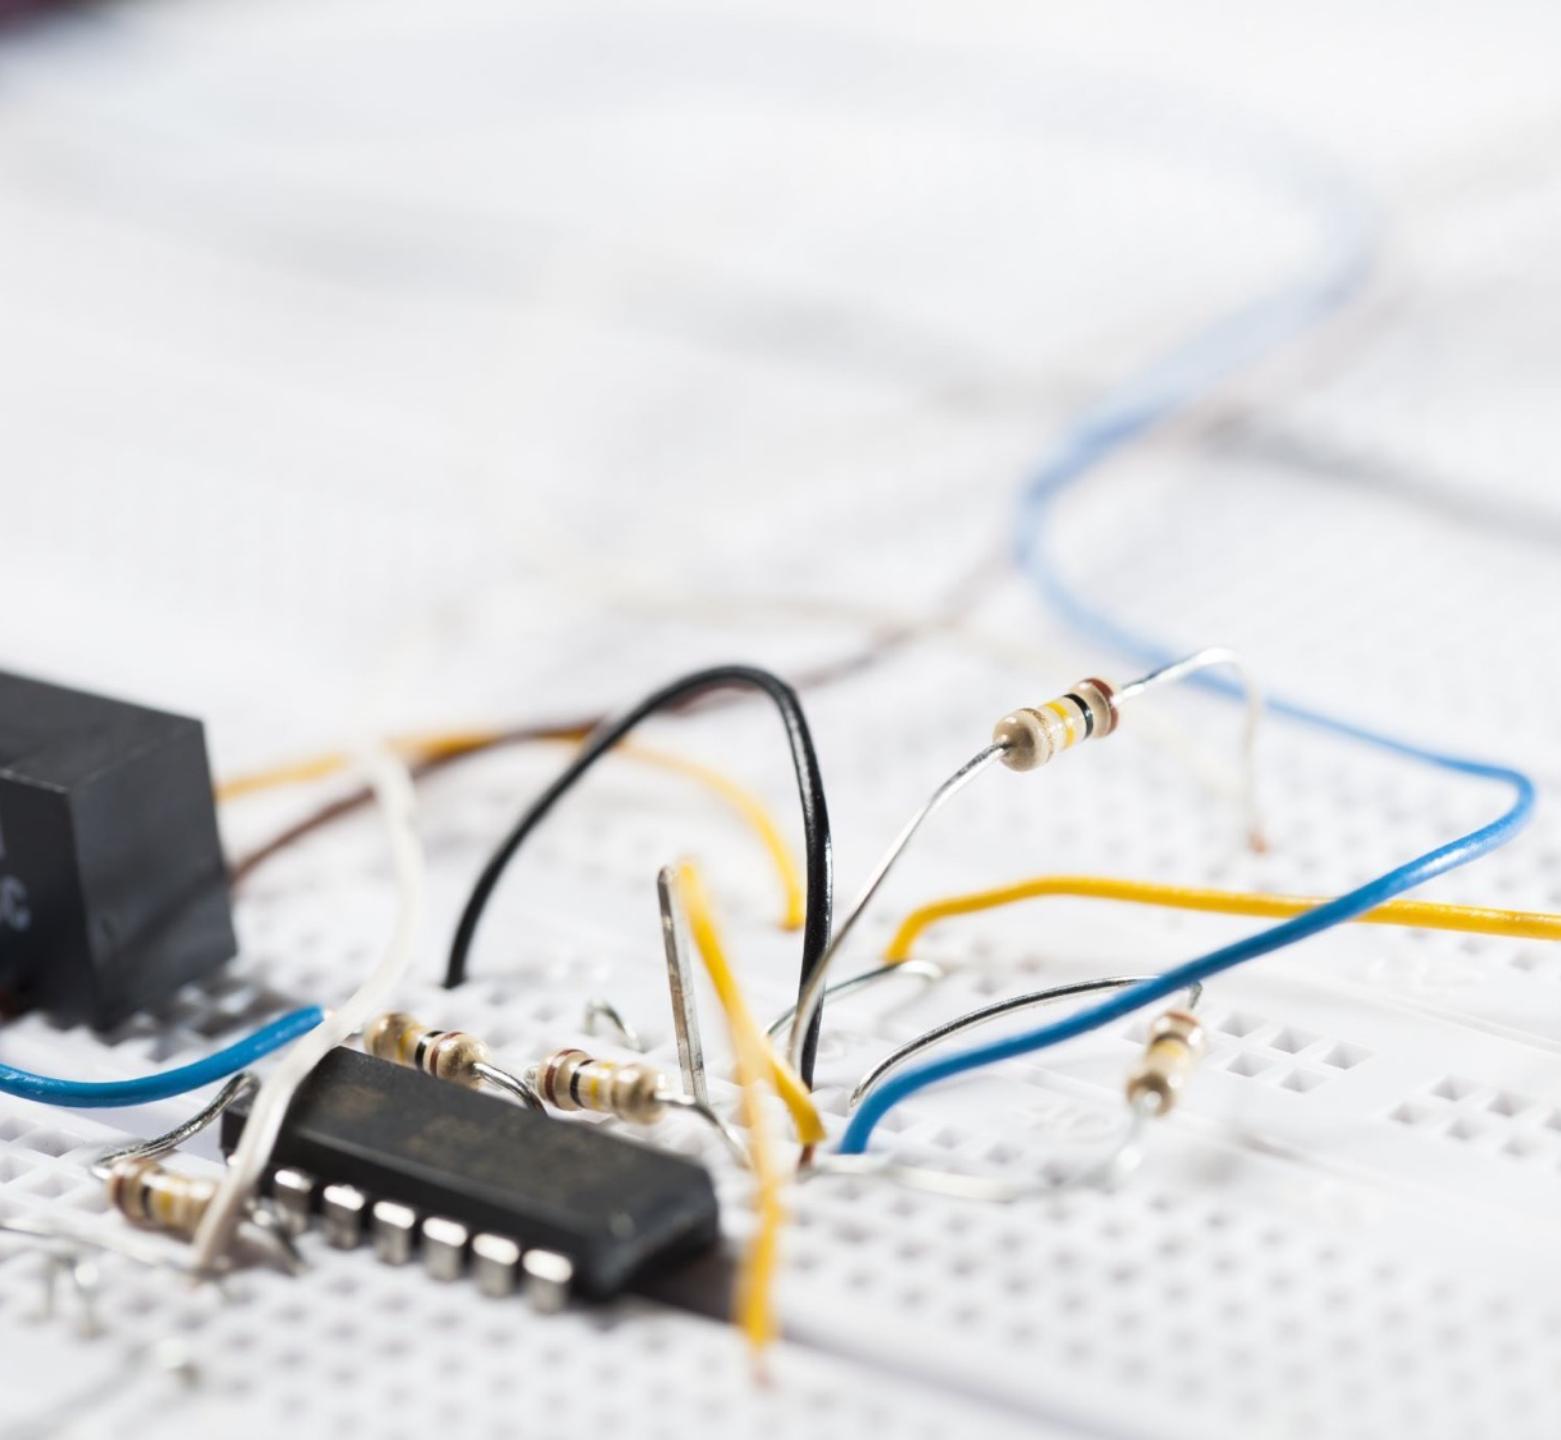
- A group of one or more containers
- If there are more than one container in a Pod, it's called a “sidecar container”
- Smallest deployable unit





# Deployment

- Manages a set of Pods
- Is considered a "high level Controller"
- Has the ability for self healing
- Deployments do not manage the Replicas/Pods, ReplicaSets do



## DaemonSet

- Like the Deployment Controller, it's a high-level Controller
- Default behavior is to ensure that a Pod is running on each Worker Node
- Purposes include any time that you need a Pod to run on each worker node. For example, if there's an agent running that's needed to run on each node
- You can specify which nodes you want the Pods to run on with Labels.



- A higher-level Controller like the Deployment and DaemonSet
- Ensures that a Pod is unique with its own unique identifier
- Maintains the sticky identity for each Pod
- StatefulSets are used for apps that require unique network identifiers, stable persistent storage, graceful deployments/scaling, and rolling updates in an ordered fashion

## StatefulSet

# Custom Resource Definition (CRD)

Extends the  
Kubernetes API

Allows you to have  
your very own  
Objects/Kind

A config that's not  
enabled by default

You can create  
your own or use  
what already  
exists

# Controllers

---

Ensures that the current state is the desired state

3 replicas? That means you should have 3 replicas

Reconciliation loop

Most k8s objects contain a Controller

A large orange circle is positioned on the left side of the slide, overlapping the white background. It is centered vertically and has a radius extending towards the bottom-left corner.

# Operators

---

Controllers + CRD's ==  
Operators

---

Kubebuilder

---

Operator SDK Framework

---

# Authentication and authorization

Authentication ==  
logging in.

Authorization ==  
permissions/RBAC

# RBAC Verb List

---

**Get**

View individual resources like a Pod

**List**

Retrieve a list of resources. Example – view all Deployments in a Namespace

**Watch**

Watch changes to resources in real-time

**Create**

Create a new resource

**Update**

Modify an existing resource

**Patch**

Make partial updates to a resource

**Delete**

Remove resources

**DeleteCollection**

Delete a collection of resources

# Backups

**Etcd**

**volumes**

# Options

Velero

Veeam

S3  
Backups

Azure  
Backups

## Core API Group vs Named API Group

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: projectx_policy
  namespace: default
```

```
apiVersion: v1
kind: Pod
metadata:
  name: static-web
```

# Declarative

---

- Tell me what to do, not HOW to do it

# Imperative

---

- Tell me what to do and HOW to do it

# Imperative and Declarative

---

```
spec:  
  containers:  
    - name: nginxdeployment  
      image: nginx:latest  
      resources:  
        requests:  
          memory: "64Mi"  
          cpu: "250m"  
        limits:  
          memory: "128Mi"  
      ports:  
        - containerPort: 80
```

## Declarative

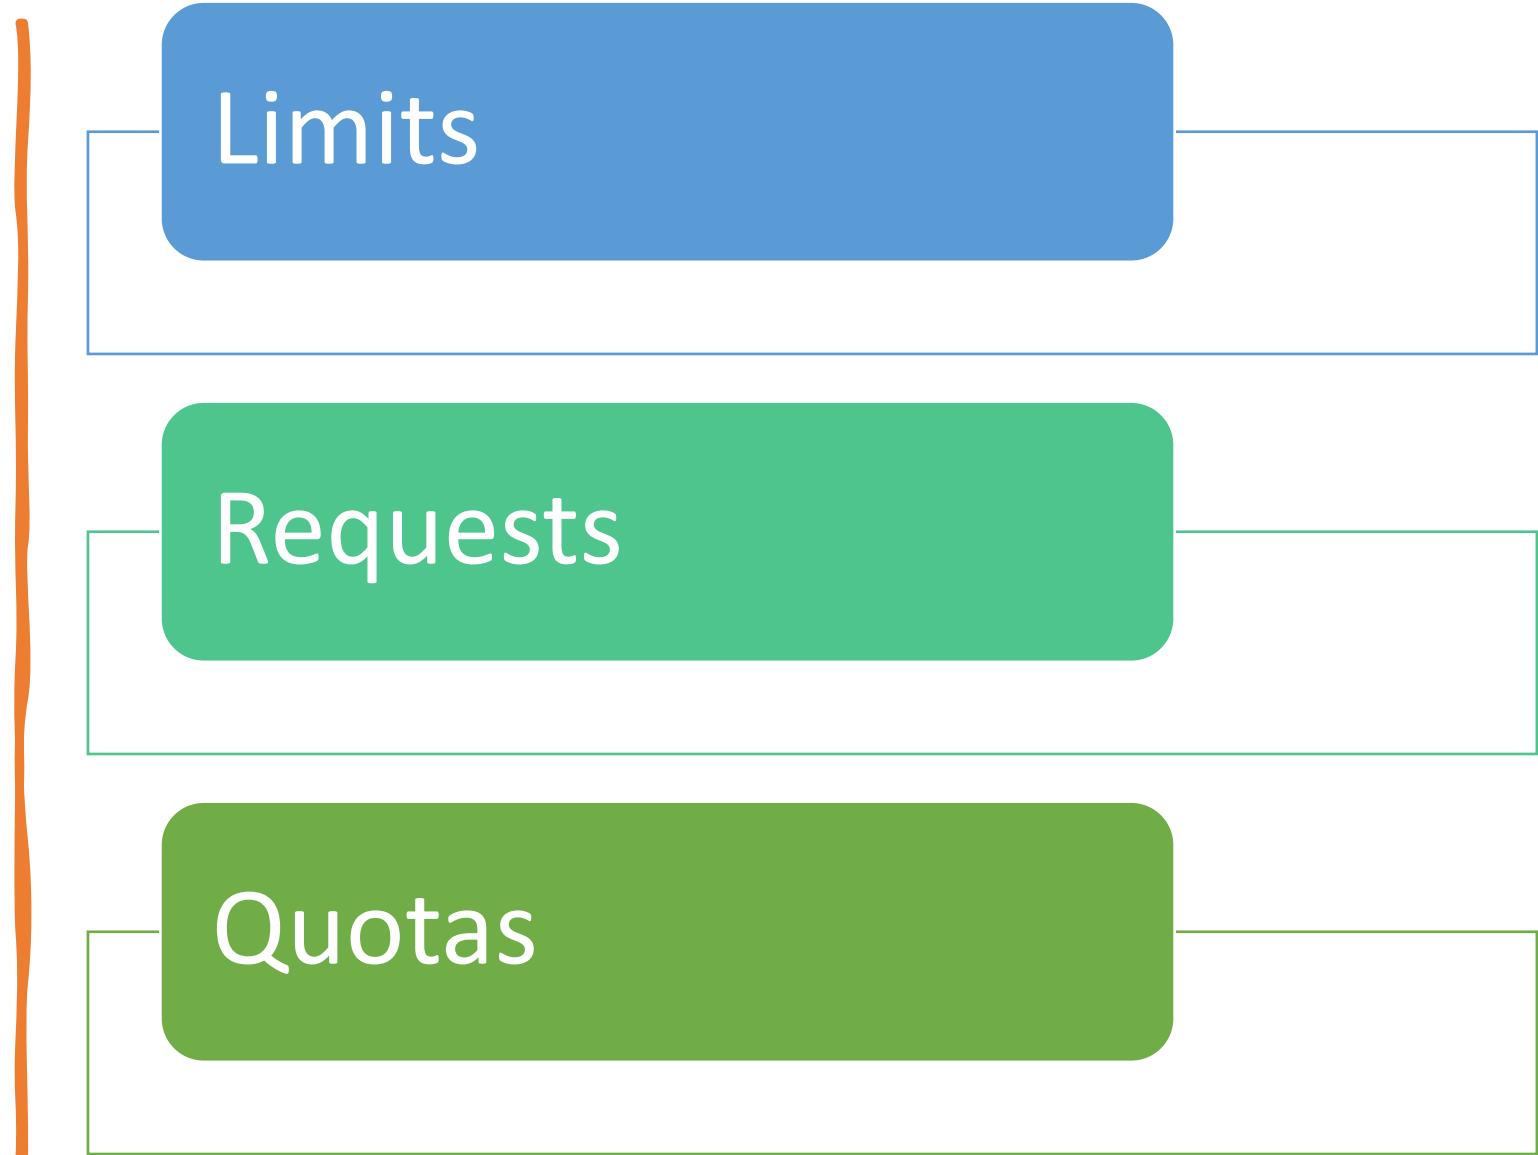
- A list of things you want done
- Code that focuses on specifying the result of what you want

```
OS=xUbuntu_20.04  
VERSION=1.22  
  
echo "deb https://download.opensuse.org/repositories/debian/"$VERSION/crio/"$OS" > /etc/apt/sources.list.d/crio.list  
echo "deb http://download.opensuse.org/repositories/debian/"$VERSION/crio/"$OS" / > /etc/apt/sources.list.d/crio.list  
  
curl -L https://download.opensuse.org/repositories/debian/"$VERSION/crio/"$OS/ | tee /etc/apt/sources.list.d/crio.list  
curl -L https://download.opensuse.org/repositories/debian/"$VERSION/crio/"$OS/ | tee /etc/apt/sources.list.d/crio.list  
  
exit  
  
sudo apt update -y  
sudo apt install cri-o cri-o-runc -y  
  
sudo systemctl daemon-reload  
sudo systemctl enable crio --now
```

## Imperative

- A list of things you want AND how to do them
- Focused on an explicit sequence of commands to describe how you want a computer to do things

# Resource Utilization



# Namespaces

## Namespaces

In Kubernetes, *namespaces* provides a mechanism for isolating groups of resources within a single cluster. Names of resources need to be unique within a namespace, but not across namespaces. Namespace-based scoping is applicable only for namespaced objects (e.g. *Deployments*, *Services*, etc) and not for cluster-wide objects (e.g. *StorageClass*, *Nodes*, *PersistentVolumes*, etc).

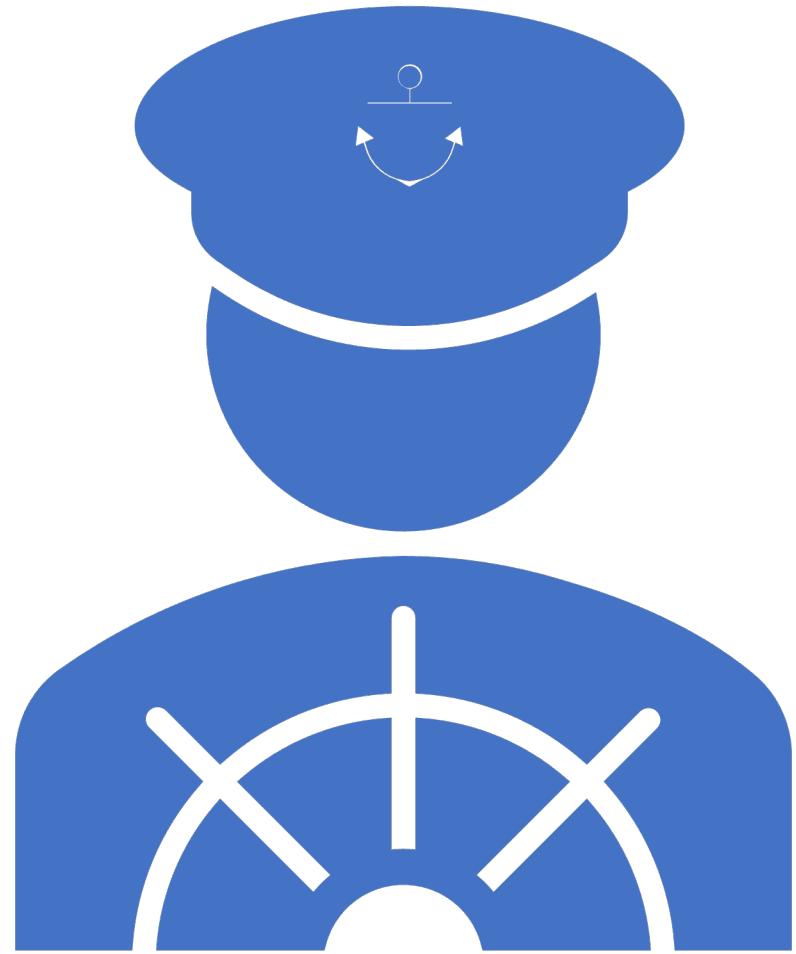
# Secrets

## Secrets

A Secret is an object that contains a small amount of sensitive data such as a password, a token, or a key. Such information might otherwise be put in a Pod specification or in a container image. Using a Secret means that you don't need to include confidential data in your application code.

Because Secrets can be created independently of the Pods that use them, there is less risk of the Secret (and its data) being exposed during the workflow of creating, viewing, and editing Pods. Kubernetes, and applications that run in your cluster, can also take additional precautions with Secrets, such as avoiding writing sensitive data to nonvolatile storage.

Secrets are similar to [ConfigMaps](#) but are specifically intended to hold confidential data.



# Helm

---

- Package Manager (think Aptitude)
- Package up Kubernetes Manifests

# Workloads and Scheduling

1

Deployments

2

Upgrades

3

Resource  
optimization and  
scaling

# What To Log



AUTHENTICATION



CONTROL PLANE LOGS

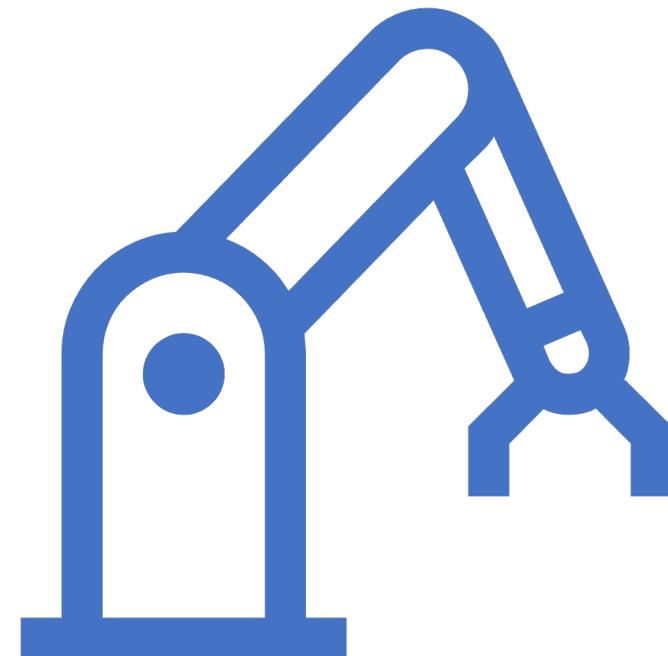
# What To Log For Worker Nodes

Network logs  
from kube-proxy,  
eBPF, or CNI

Kubelet

# Logging Options For Pods

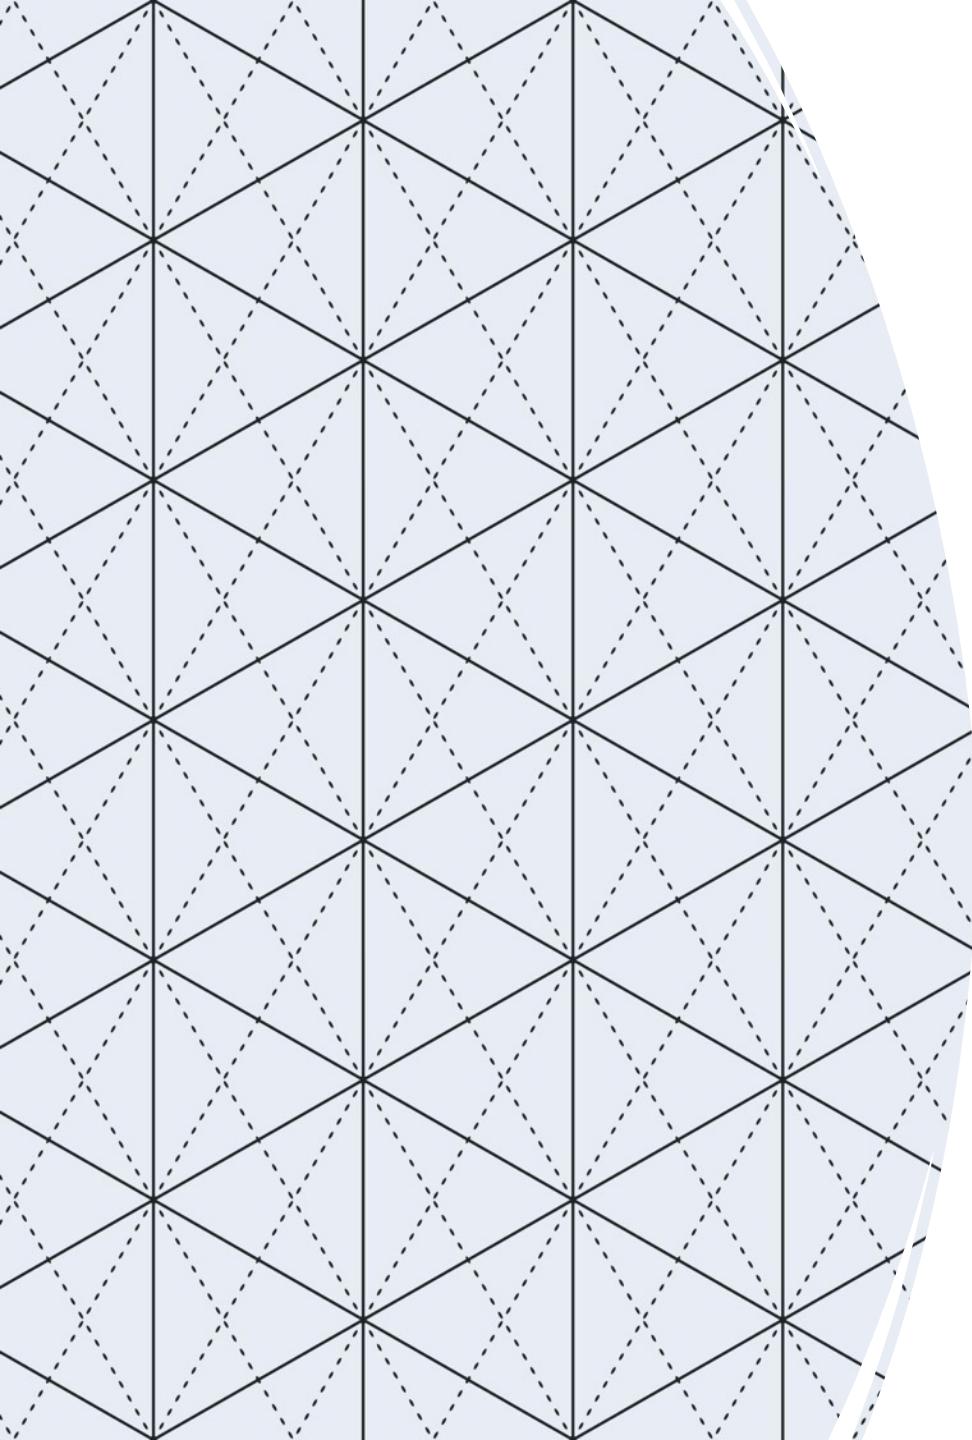
- Stdout and Stderr
- Using Application-level Logging Configuration
- Using Logging Agents



# Commands

```
kubectl logs  
pod_name
```

```
kubectl describe  
*resource*  
resource_name
```



## Monitoring

- Seeing the data in real-time
  - Having the ability to go back 30/60/90/Etc. days to see what happened
  - Alerting capabilities
-

# Observability

- Doing something with the data

## Metrics

Collecting time series data, which is used to predict expected ranges and forecast values, showing it in dashboards (like Grafana or another UI-centric dashboard), and alerting on it.

## Traces

The health of an application from an end-to-end perspective. It's the process of investigating how events take place between microservices.

# Logs

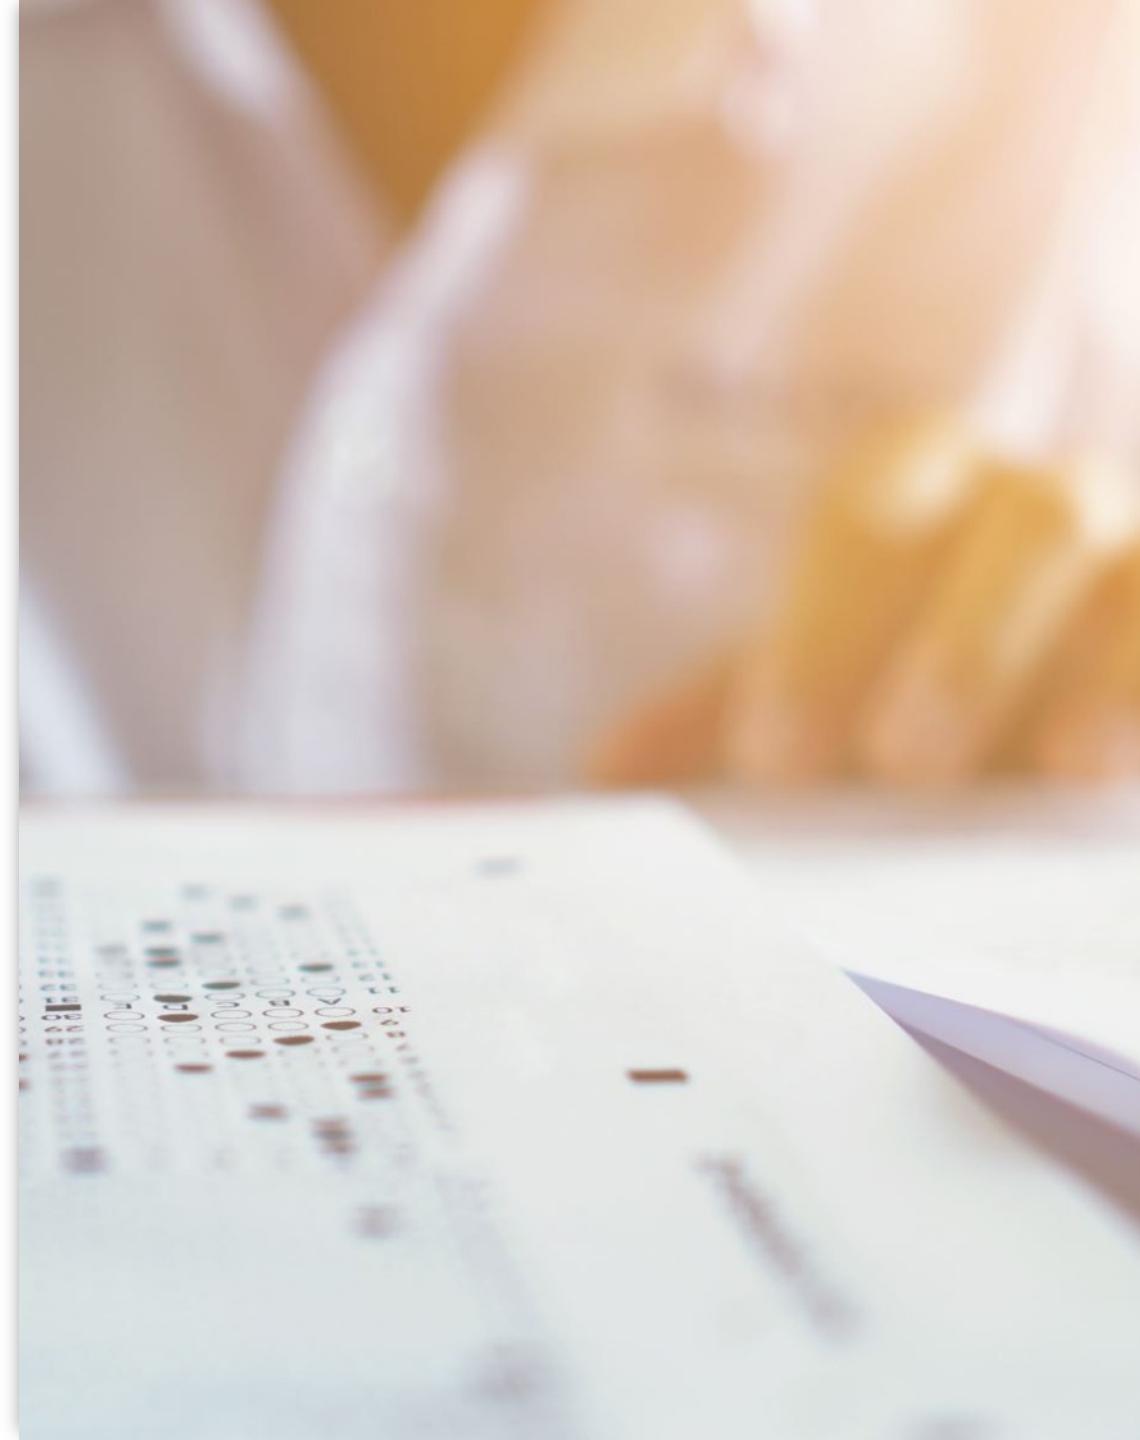
- Aggregating and  
storing logged  
event messages  
written by  
programs and  
systems

---



# Policy Enforcement

- OPA
- Kyverno
- Admissions Controllers



# OPA

Gatekeeper

# OPA Implementation

01

Config = Definition  
of what Gatekeeper  
is allowed to create  
policies for

02

Constraint Template  
= Rule/policy that  
you want to  
configure for your  
environment

03

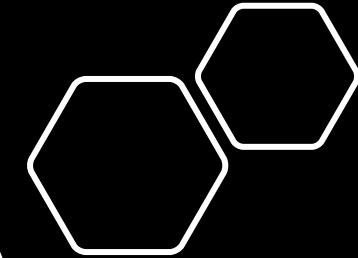
Constraint Config =  
Bringing the  
constraint to life  
(implementing it)

# Pod vs Container Logging

Kubernetes captures logs from each container within a Pod

You can see logs per Pod

Those logs can be (and should be) exported

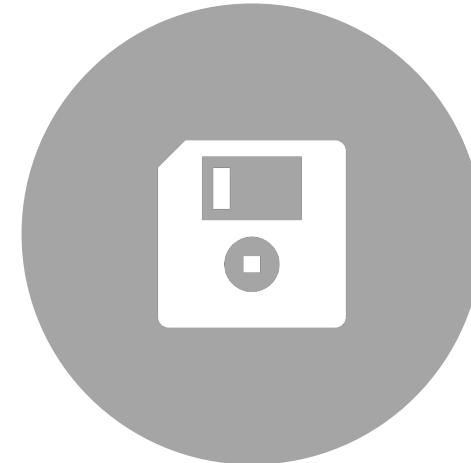


# Etcd Backup

---



IN THE CLOUD, THERE ARE  
SPECIFIC SERVICES



ON PREM, YOU WOULD HAVE TO  
CHOOSE A BACKUP SOLUTION

# Troubleshooting On-Prem and Cloud Networks

---

1

Ensure that the host networking is working as expected

2

Ensure kube-proxy (or eBPF) is working

3

Ensure the CNI is working

# Container Logging

# Logging

---

Stderr

1

Typically used to output error messages

Stdoutt

2

Command normal output

# CSI

A way to use storage (like a hard drive) on  
Kubernetes



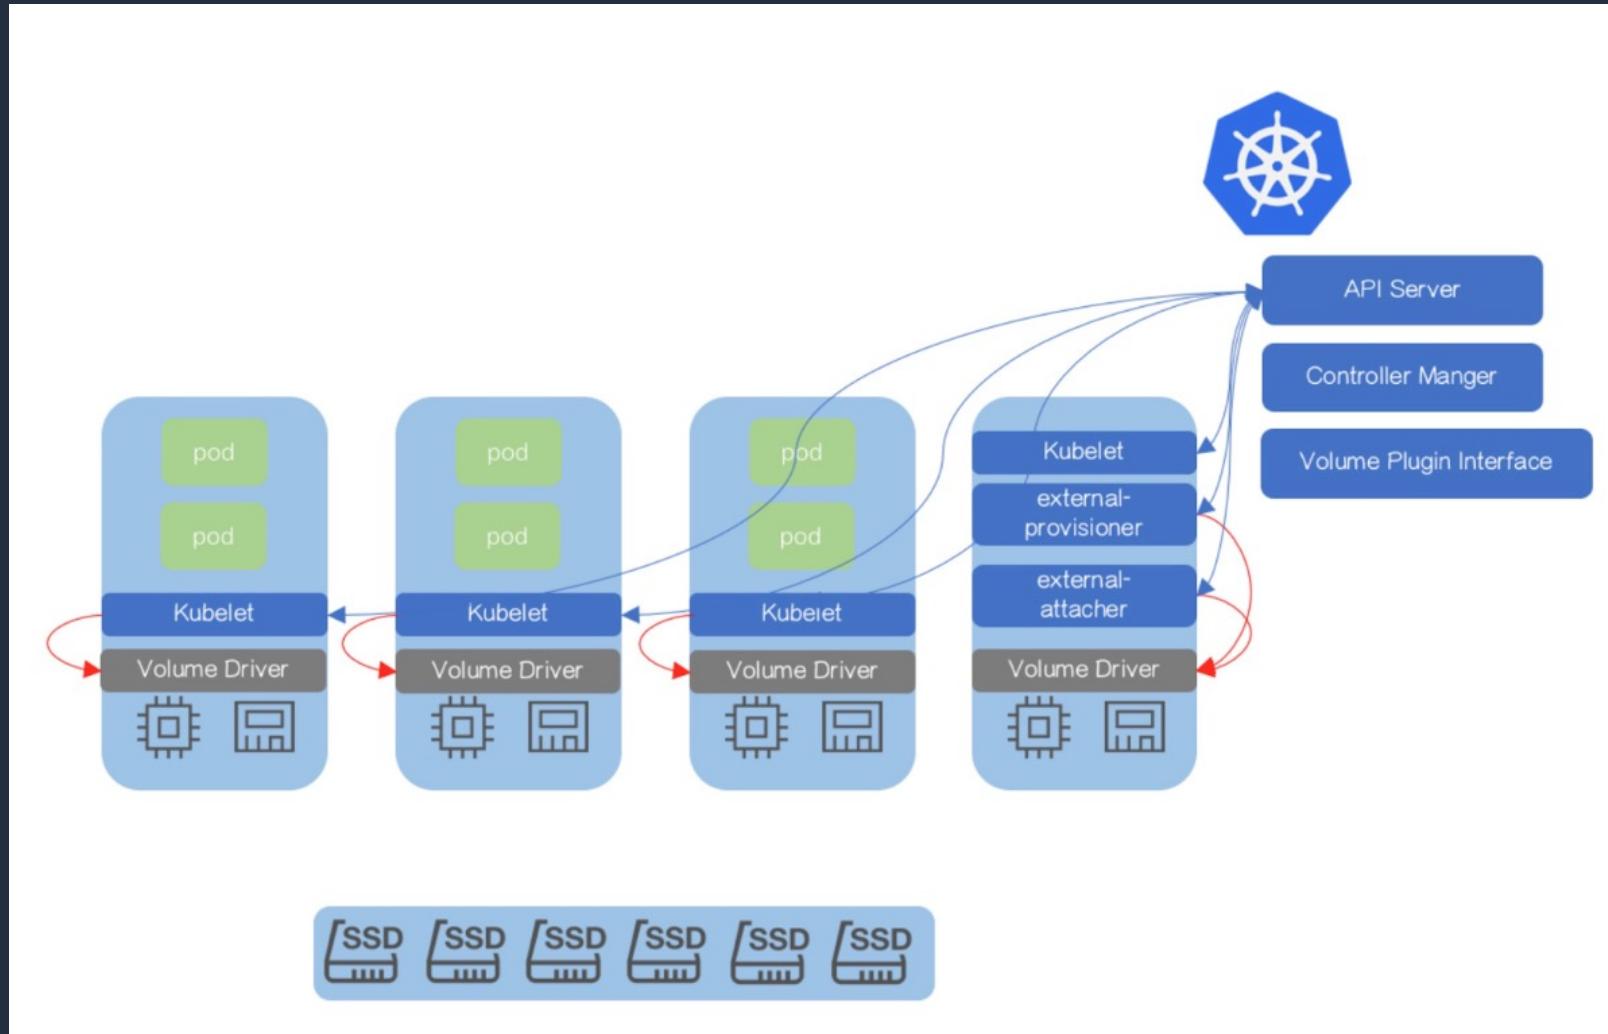
---

---

# CSI Breakdown

- Gives Kubernetes the ability to use storage, which is called Volumes.
- Various storage interfaces available
- It's a plugin much like the CNI and CRI

# CSI Breakdown



# Volume Reclaim Policies



# Retain

---

---

- If a user deletes a PersistentVolumeClaim, the corresponding PersistentVolume will not be deleted
- It is moved to the Released phase, where all of its data can be manually recovered.

# Recycle

- This policy is deprecated
- **Performs a scrub on the volume and makes it available for re-use**

# Delete

---

---

- Default action for dynamically provisioned PersistentVolumes
- **A dynamically provisioned volume is automatically deleted when a user deletes the corresponding PersistentVolumeClaim**
- This automatic behavior might be no good if the volume contains data that you need

# Kubernetes Secrets

---

- Stored in Etcd
- **Stored as Base64 Plain Text**

# What You Should Know

1. Linux
2. Networking
3. Storage
4. Lite programming/automation
5. Application support
6. Infrastructure/scaling

# Home Lab With Cash

Synology DS223 Diskstation NAS  
Intel NUC 12 Extreme Kit NUC12DCMi7

# Home Lab With Less Cash

Intel NUC 9 Business Mini PC Desktop  
Synology DS120j 1 bay NAS

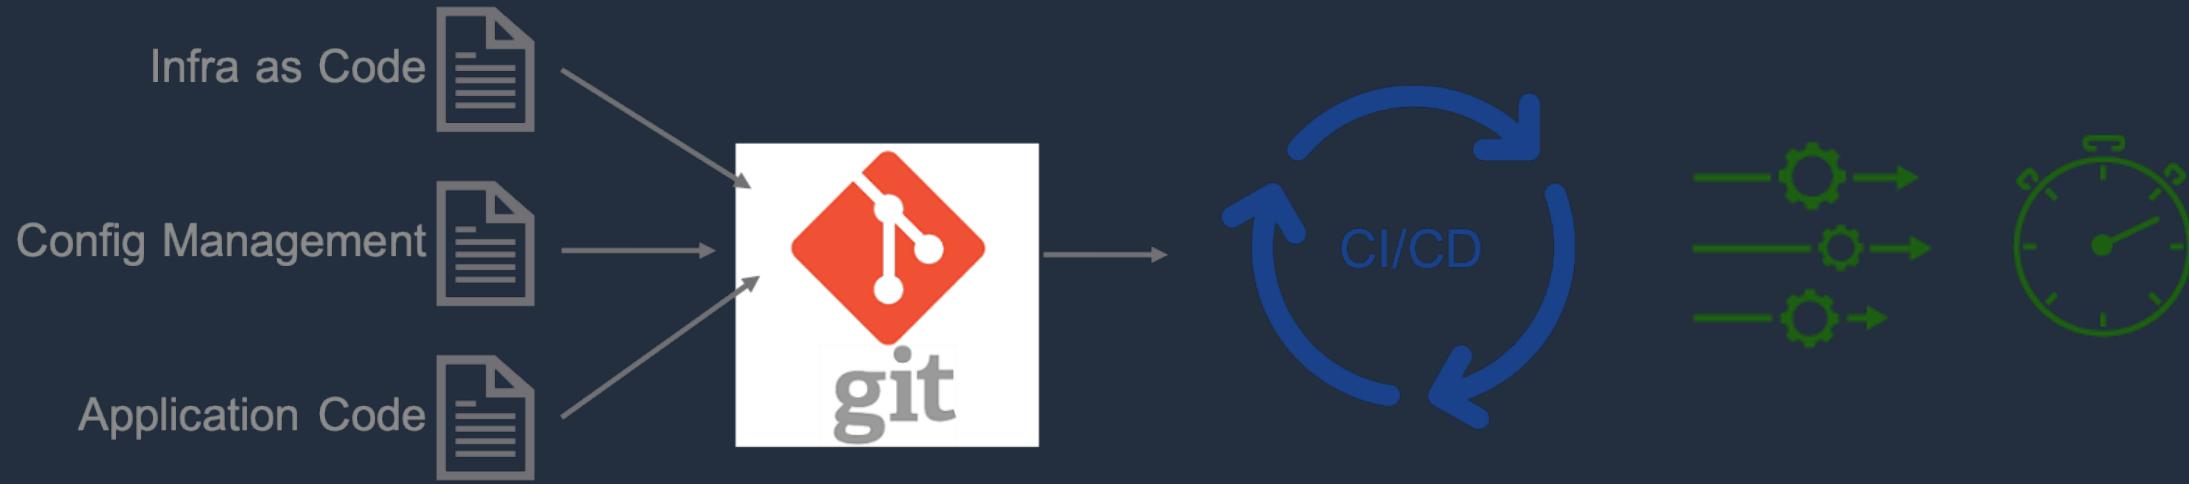
# GitOps

---

---

- Deploy our code automatically
- Confirm current state is the desired state
- Git-based (like the name suggests)

# GitOps-in-a-nutshell



“Source of Truth” for declarative code

Update to code source triggers a pipeline

Pipeline runs a series of tasks, resulting in the update of the runtime environment to match the source

---

---

# Tools

- Have there been any tools requested by teams?
- Do you know the business value of what you're implementing?
- Can the engineers in-house get familiar with the tools in a timely fashion?
- What impact would bringing on a new tool create?

# Container Network Interface (CNI)

---

- Dynamically configure network resources
- A framework/plugin to make Networking a separate entity for Kubernetes resources

# Cluster Networking

Outside of the CNI

# Kubernetes Service

---

- “frontend” of Pods
- A way to access Pods in a more production-level fashion
- Definition: A Service is a method for exposing a network application that is running as one or more Pods in your cluster
- Services get DNS Names, which come from CoreDNS

# LoadBalancer vs ClusterIP vs NodePort

1

LoadBalancer

Creates a service that has a load balancer attached to it

2

ClusterIP

Default type. Used to create a connection using the clusters IP and a port. Used for communication within the cluster

3

NodePort

Exposes a port on each node. Like a ClusterIP, but used for communication within and outside of the cluster

# Network Policies

---

---

- Firewall rules for Pods
- Controls Ingress and Egress traffic

If you want to control traffic flow at the IP address or port level (OSI layer 3 or 4), then you might consider using Kubernetes NetworkPolicies for particular applications in your cluster. NetworkPolicies are an application-centric construct which allow you to specify how a pod is allowed to communicate with various network "entities" (we use the word "entity" here to avoid overloading the more common terms such as "endpoints" and "services", which have specific Kubernetes connotations) over the network. NetworkPolicies apply to a connection with a pod on one or both ends, and are not relevant to other connections.

# Ingress and Egress

Traffic in

Traffic out

# Kubernetes Controllers

Confirm that the current state is the desired state

# Ingress Controller

